

# Ansh Agnihotri

+91 9354616027 — [agnihotriansh05@gmail.com](mailto:agnihotriansh05@gmail.com) — [linkedin.com/in/ansh-agnihotri](https://linkedin.com/in/ansh-agnihotri) — [github.com/agnihotri-anxh](https://github.com/agnihotri-anxh)

## Education

### IIIT Senapati Manipur

*B.Tech in Electronics & Communication Engineering*

Dec 2022 – Present

CGPA: 7.65/10

- Relevant Coursework: **Data Structures & Algorithms (DSA)**, **Object-Oriented Programming (OOP)**, **Operating Systems**, **Computer Networks**, Database Management Systems (DBMS)

## Experience

### Taxmann

*AI NLP Intern*

Jan 2026 – Present

New Delhi, India

- Engineered robust **Python backend parsing utilities** to clean, structure, and ingest massive, unformatted internal datasets into modern LLM orchestrators.
- Developing an enterprise-grade **RAG-based legal search engine** using **Milvus (Zilliz)** to enable scalable semantic queries over millions of high-dimensional vector embeddings.
- **Automated** IBC Book amendment processing by utilizing **Claude** and precise **Prompt Engineering** to parse and format legal updates for direct integration into **Adobe InDesign**.

### Zidio Development

*Software Engineer Intern*

May 2025 – June 2025

Remote

- Developed and documented clean **REST APIs** using Python and FastAPI to manage user profile and authentication workflows for a SaaS dashboard.
- Assisted in writing modular backend service logic, helping improve overall system data validation and reducing API response errors.

## Projects

### HealthMate | Flask, LangChain, Groq API, Pinecone, Docker, AWS ECR, GitHub Actions [\[GitHub\]](#)

June 2025

- Developed a scalable medical AI application utilizing **Groq Llama 3.1** via a custom-prompted RAG architecture, managing document indexing across vector databases with **95%** precision.
- Containerized application services using **Docker** and established a secure **CI/CD pipeline** via GitHub Actions for seamless production deployment to AWS.

### AI-Powered Mock Interviewer | Flask, FAISS, MongoDB, EdgeTTS [\[Demo\]](#)

July 2025

- Built an asynchronous, real-time voice synthesis platform, engineering the backend to optimize context windows and achieve **sub-second (0.8s)** end-to-end user-to-agent latency.
- Designed systemic prompt flows combined with a **FAISS vector index** to dynamically synthesize contextual questions, preventing application looping and repetitive states.

### Loan Default Prediction Platform | Python, FastAPI, Gradient Boosting, PostgreSQL [\[Code\]](#)

Aug 2024 – Nov 2024

- Designed and built clean **FastAPI web services** to expose predictive ML algorithms, utilizing stratified processing to efficiently handle severely imbalanced datasets.
- Programmed modular feature engineering code blocks to isolate data transformation scopes, eliminating target leakage during production runs.

## Technical Skills

- **Languages:** Python (NumPy, Pandas), C++, SQL, HTML/CSS, JavaScript, React, Node Js
- **AI Engineering & Prompting:** Prompt Engineering, LLM Orchestration, RAG, OpenAI API, LangChain, Embedding Optimization
- **Backend & DevOps:** Flask, FastAPI, RESTful APIs, Docker, Git, GitHub Actions, AWS (ECR, EC2)
- **Databases & Vector Search:** Milvus (Zilliz), MongoDB, Pinecone, FAISS, PostgreSQL
- **AI Developer Tools:** GitHub Copilot, Cursor, ChatGPT, Claude

## Certifications & Achievements

Tata Group GenAI Powered Data Analytics Job Simulation (Forage) [\[Link\]](#)

Top 5% (Gold) in NPTEL Fundamentals of AI (IIT Madras) — Elite (Top 10%) in NPTEL Deep Learning

Top 12% Globally in Kaggle Backpack Price Prediction competition (outperformed 3,200+ participants).